

Ad Hoc Networking


Main source

Ad Hoc Networking, C. E. Perkins, Addison-Wesley

What is an ad hoc network?


- A **short lived network** just for the communication needs of the moment
- Infrastructure-less network
- Commercial Applications
 - Emergency services
 - Damaged infrastructure
 - Difficult to set up infrastructure
 - **Vehicular networks**

Model of operation



A node in an ad hoc network

Two main components of a node



Routing methods

- Recall: 2 routing algorithms on fixed networks
 - Distance-Vector (DV): Routing Information Protocol (RIP)
 - Link State: Open Shortest Path First (OSPF)
- Destination-Sequenced Distance Vector (DSDV)
- Dynamic Source Routing (DSR)
- Location Aided Routing (LAR)
- Ad-hoc On-demand Distance Vector (AODV)

Two routing algorithms for fixed networks


Distance Vector (RIP: Routing Info. Protocol)

Link State (OSPF: Open Shortest Path First)

Distance Vector

– Based on: Distributed Bellman-Ford algorithm

- Node **i** maintains, for each dest **x**, a set of distances $\{d_{ij}(x)\}$, where **j** is a neighbor of **i**.
- Node **i** treats neighbor **k** as a next hop for a packet for **x** if $d_{ik}(x) = \min_j \{d_{ij}(x)\}$.
- To keep $\{d_{ij}(x)\}$ up to date,
 - Node **i** monitors the **cost** of its outgoing links.
 - Periodically broadcast your estimate of the shortest distance to every other node.




Distance Vector (contd.)

- Advantages: efficient, easier to implement, less storage.
- Drawback: short/long-lived loops.
 - Two-node and three-node instability problems
- Ad hoc networks → rapid topological changes → loops.

Link State

- Each node has its own **view** of the net topology and **link costs**.



- For **consistent view**
 - Each node periodically **broadcasts** outgoing link **costs**.
 - Based on received info, nodes **update** their **view** of net.
- Apply **shortest-path algorithm** to **choose** its **next hop** for each destination.

*DS*DV Routing Algorithm

(Enhancement of Distance Vector routing)

DSDV Protocol Overview


- **Each node** constructs and maintains a **route table**.
- Route table entry of **node i**:

Destination (node j)	Next hop (NH)	# of hops	Sequence #

Sequence #: originated by the **destination node**.

Seq #: Sometime before, node j made a broadcast with a “seq #” that was propagated to i via NH.

DSDV – Sequence number




Overview (contd.)

- To **maintain consistency** of route tables, each node
 - Periodically transmits **updates**.
 - Transmits **immediately** when **significant** new info is available.
- Update information consists of
 - **Which nodes** are **accessible** from the node.
 - **Number of hops** necessary to reach them.
 - (Sequence #)

Route Advertisement

- Structure of a packet **broadcast** by each node

Seq #1 (new)	Dest (j)	# of hops to dest	Seq #2
-----------------	----------	----------------------	--------


...

Generated by
broadcasting node **i**

Seq #2: Sometime before, node j made a broadcast with a sequence # that was propagated to i via NH and i chose NH to reach j.

Route selection criteria



- When a node **receives info.** about **a new route to a node**
 - **Compare** it with existing info. of a route to the same node.
 - **Use** a route with a **more recent sequence number**.
 - **Use** a route with an **identical** sequence number, if it has a **better metric**.
 - » The existing route is discarded or saved as a less preferable route.
 - **Increment** the metrics for routes chosen from the newly received broadcast info by one hop.

Responding to **topology change**

- What is a topology change?
 - **Discover** a new node from a received advertisement
 - **Learn** about a broken link with a neighbor.
 - » How: No communication from the neighbor for a long time.
 - » Update the route table as follows:
 - » Assign an **infinity** metric to that destination.
 - » Assign a **new** sequence number. (last seq # received from that node + 1) **←** As if the node told us that it moved away
- What to do next because of this **substantial** change:
 - Disclose this new info in a **broadcast** packet.

Responding to topology change

- When a node **receives** an **infinity** metric and it has a “=” or “>” sequence # with a **finite** metric
 - Trigger a route update broadcast.



Efficient Response to topology change

- To **reduce** the amount of info in broadcast packets, two types of broadcasts are used:
 - **Full dump**: All available routing info., periodic
 - **Incremental**: Info changed since last full dump

Dynamic Source Routing (DSR)

Note: It is the responsibility of the source node to **find and maintain** routes.


Source Routing

- The source of a packet **determines the complete sequence of nodes** (i.e. the route) for the packet.
- The route is put in the packet's **header**.
- There is no periodic route advertisement.
- When a node needs one, it dynamically determines one based on
 - Cached information
 - Results of a route discovery protocol

Basic Operation

- To send a packet, construct a source route in the packet header.
- The sender transmits the packet to the first node on the **route**.
- When a node receives a packet, if it is not the final destination, it simply transmits the packet to the next node on the route.
- Each node maintains a **route cache** in which it caches source routes that it has **learned**.
- When a source wants to send a packet to another node, it first checks its route cache:
 - If a route is found, use the route.
 - If no route is found, discover one.

Learning routes from others



Route Discovery (find a route)

- A node initiating a route discovery **broadcasts** a **route request** (RR) packet (to all its neighbors).
 - RR header: Sequence#, source, destination.
- A node receiving an RR *is NOT* the destination:
 - The node further broadcasts it after inserting its ID in the RR header. ➔ Incrementally construct a route.
- A node receiving an RR *is* the destination:
 - It sends a **route reply** packet containing the route (sequence of nodes from the source to the destination).
 - A route reply takes the reverse of the defined path just discovered by the RR. ⬅ Assumption: Symmetric communication between nodes.


Route Discovery (efficiency)

- Each node maintains a **list** of
 <source ID, RRseq. #> for **all** sources
- When a node receives a route request, it does this:
 - If the pair <source ID, RRseq. #> from the RR is already on the above list, **discard** the packet.
 - If the node's address is already in the route record in the request, **discard** the packet.
 - If the target of the request matches this node's address, **send a route reply**.
 - Else, insert this node's own address in the route record, and **rebroadcast** the packet.

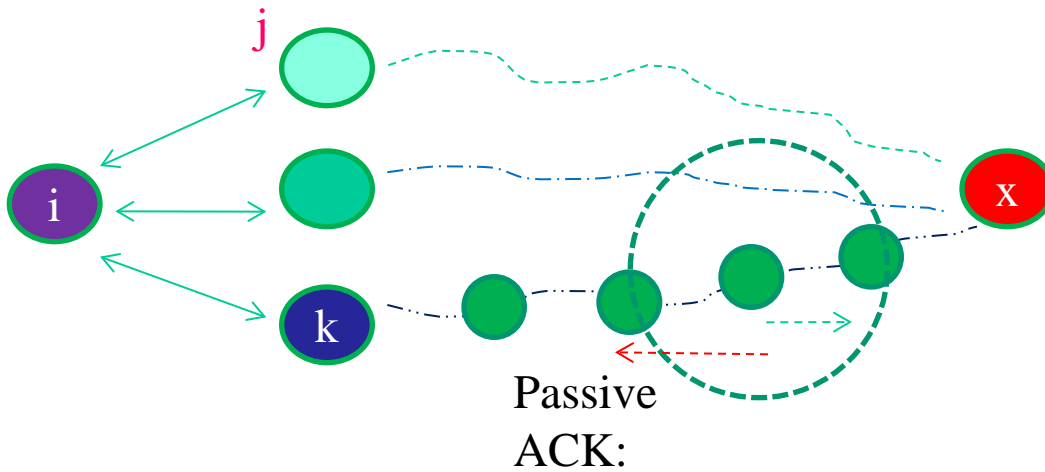
Route Maintenance

- While a route is in use, a route maintenance procedure **monitors** the operation of the route and **informs** the **sender** of any routing error.
 - Wireless networks generally utilize hop-by-hop ACK. If the data-link layer in a node on a route reports a problem, the node sends a route error packet to the original sender.
 - If link-level error detection is not available, you may use the idea of passive ACK (i.e. being able to hear that node transmitting the packet again.)
 - Worst case: Ask for an explicit ACK.

Route maintenance



Route maintenance



Location Aided Routing (LAR)

1. **Obtain** location information from a GPS receiver
2. **Utilize** location information to improve the performance of routing protocols

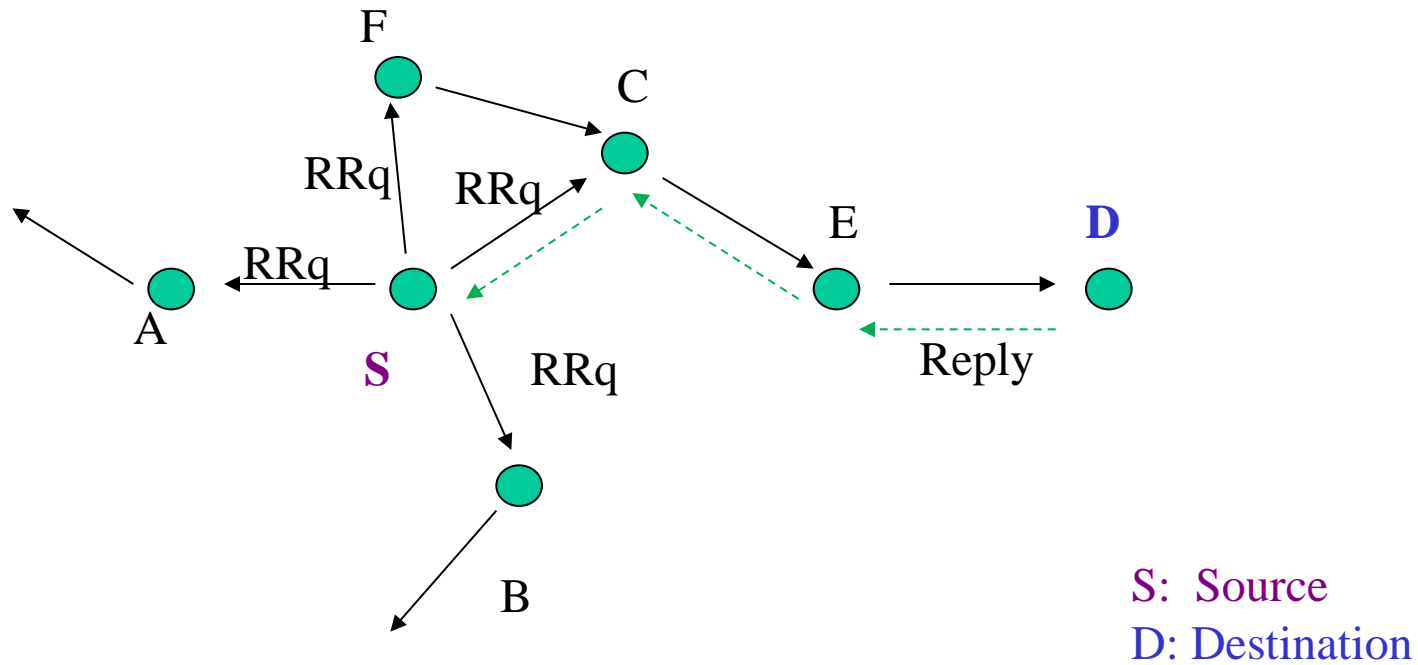
Motivation

- Host mobility causes topology change
 - ➔ Task of **discovering** and **maintaining** routes is non-trivial
- Flooding (ex. DSR) is a brute force way of discovering routes.
- LAR: **Reduces the full impact of flooding** by forwarding route discovery messages to a selected region by using location information.

Route discovery using flooding

- A sender broadcasts a route request (RRq) to all its neighbors
- If a node *is NOT* the destination of an RRq, it broadcasts the RRq.
 - Subsequent copies of an RRq are not rebroadcast: save recent <Sender, sequence #> pairs.
- If a node *is* the destination of an RRq, it *sends a reply* back. The reply takes the reverse path taken by the RRq.
- **Timeout:** If the sender does not receive a reply within a certain timeout period, it reinitiates the process.

Route discovery using flooding



LAR: Preliminaries


- Assumptions
 - Node S needs to find a route to D.
 - S knew D's location "L" at time t_0 (in the past)
 - Current time is t_1 .
- **Expected Zone** of D from the viewpoint of S at t_1
 - This is the region that S expects to contain D at time t_1 .
 - Zero knowledge: the entire network area is the expected zone.
 - Having more information regarding mobility of a destination node can result in a smaller expected zone.

LAR: Preliminaries

- **Request Zone**

- Node S defines a request zone for the route req.
- To increase the probability that the route request will reach node D, the request zone includes the expected zone.
- Additional regions must be included in the request zone so that S and D belong to the request zone.
- A node **forwards** a route request only if it lies in the request zone. ← Cost reduction
- If a route is not discovered within a suitable time period, S reinitiates route discovery with an EXPANDED request zone.

Request Zone



Tradeoff between probability of finding a route and discovery overhead.

The LAR Algorithm


- It is the FLOODING algorithm.
- A node that is NOT in the request zone does not forward a route request to its neighbors.

Two ways of computing a Request Zone

- LAR Scheme 1
- LAR Scheme 2

LAR Scheme 1

- A request zone is the smallest rectangle that includes the current location of S and the expected zone of D, such that the sides of the rectangle are parallel to the X- and Y-axis.
- The source explicitly specifies the request zone in its route request.





When D receives a Route request, it sends a route reply which includes its **current location** and the **current time**.

LAR Scheme 2

- Sender S includes these info with a route request.
 - $DIST_S$: Distance of S from D
 - (X_d, Y_d) : Coordinates of D
- When a node I receives the route request from S, I calculates its distance $DIST_I$ from D and acts as follows:
 - If $DIST_S + \delta \geq DIST_I$ I forwards the request to its neighbors after replacing $DIST_S$ with $DIST_I$.
 - If $DIST_S + \delta < DIST_I$ I discards the request.
- Idea: A node (I) forwards a request forwarded by another node (S), if I is “at most δ farther”.

Request Zone




AODV: **A**D HOC **O**N-DEMAND **D**V ROUTING PROTOCOL

AODV: Ad Hoc On-demand DV


- Does **not** attempt to maintain routes from every node to every other node.
- Routes are discovered on an **as-needed basis** and are maintained as needed.
- The protocol uses three kinds of messages:
 - **RREQ**, **RREP**, **RERR**
- **Need:** A node wants to send a message to a destination
 - If a route exists in its RT to that dest., forward the pkt to the **next hop**; otherwise, start a **route discovery process**.

AODV (Example)




→ RREQ

AODV (Example)




-----> Reverse
Path Setup


AODV (Example)



AODV (Example)




AODV (Example)




→ RREP

AODV (Example)




Forward
Path Setup


AODV (Example)



AODV (Example)



AODV (Example)



AODV: Route Discovery Process

- Create a **RREQ**(BcastID, SrcIP, CurSeq#, DestIP, LastKnownSeq#)
 - BcastID is incremented for each **new** discovery.
 - (BcastID, SrcIP) uniquely identifies a discovery process.
 - **Note:** Every time a node sends out any type of message, it increases its own sequence number: **CurSeq#**.
- Broadcast **RREQ** and start a timer.
- When a node receives a **RREQ**, it checks if it has **seen** the **RREQ** before. If YES, discard the **RREQ**, else **process** it as follows.
- Processing a **RREQ** (see the next slide.)


AODV: Route Discovery Process (Processing RREQ)

- The node sets up a **reverse route** entry for the source node in its RT. [Reverse Route: SrcIP, SrcCurSeq#, # of hops to Src, IP of the neighbor that gave it RREQ, Lifetime]
- To respond to the RREQ: **2 conditions** must hold.
 - The node must have an **unexpired entry** for the Dest.
 - The **Seq#** associated with that Dest must be **at least as great as** that indicated in the RREQ.
- (The above conditions **prevent the formation of loops** by ensuring that the route returned is never old enough.
- If the node **satisfies** the above 2 conditions, it responds by unicasting a **RREP** back to the source (See 6.2.3.)
- If the node is **unable to satisfy** the RREQ, it increments the RREQ's hop count and then broadcasts the RREQ to neighbors.
- **Naturally**, the Dest is always able to respond, eventually.

AODV: Route Discovery Process

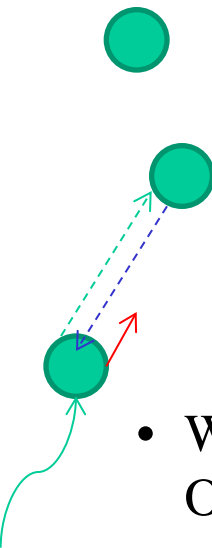
- **Forward Path Setup**

- **Current Enough Route:** Any route with a Seq# not smaller than that indicated in the **RREQ**.
- When a node determines that it has a route **current enough** to respond to the **RREQ**, it creates a **RREP**.
 - » **If the Dest** is responding, it places its **CurSeq#** in RREP, initializes **HopCount = 0**, and initializes the **LifeTime** field of the route.
 - » **If an intermediate node** is responding, it places its record of the DestSeq# in the pkt, sets the HopCount to its distance to Dest, and calculates the amount of time for which its RT entry will still be valid.
- It then unicasts the **RREP** toward the Src, using the node from which it received the **RREQ** as the next hop.



AODV: Route Discovery Process

- **Forward Path Setup (Contd.)**


- When an intermediate node receives the **RREP**, it sets up a **forward path** entry to the Dest in its RT. The entry contains **DestIP**, **NeighborIP** who gave **RREP**, **DestSeq#**, **HopCountToDest**, **LifeTime** (as in **RREP**).
A diagram showing three green circular nodes. The bottom node has a red arrow pointing up to the middle node. A dashed blue line connects the middle node to the top node. A green curved line points to the bottom node.
- Note: Each time the route is used, its LifeTime is **updated**. If it is NOT used within the specified LifeTime, it is deleted.

- When a node receives a **RREP** for a given Dest from more than ONE neighbor:
 - It forwards the **FIRST RREP** and forwards a later **RREP** only if that **RREP** contains a **larger DestSeq#** **or** a **smaller HopCount**. Otherwise discard the **RREP**.
 - A source can set multiple routes to the same Dest.

AODV: Route Discovery Process

- **Route Maintenance**

- Once a route has been discovered for a given **Src/Dest** pair, it is **maintained** as long as needed.
- **Movement of nodes** affect only the routes containing those nodes. (Such a path is called an **active path**.)
- Movement **not along an active path** triggers no protocol action.
- IF a **Src** moves during an active session, it can reinitiate route discovery to establish a new route to the **Dest**.
- WHEN a **Dest/intermediate node** moves, a Route Error (**RERR**) is sent to the affected Src nodes.
 - » This is **initiated** by the node **upstream** of the break (closer to the Src.)
 - » It **lists** each of the Dest that are now unreachable because of the loss of the link.



AODV: Route Discovery Process

- **Route Maintenance (Contd.)**

- When the neighbor(s) receive the **RERR**, they mark their routes to Dest “Invalid” by setting the distance to **INFINITY** and propagating **RERR** to their precursors. When a Src receives the **RERR**, it **can reinitiate** route discovery.

- **Local Neighborhood**

- “Hello” messages (Unsolicited RREP) with own IP and Current Seq#.
- Not rebroadcast beyond 1-hop.

AODV: Route Discovery Process

- **Actions after Reboot** (reboot: loss of Seq# info, neighbors)
 - Wait for **delete_period**, during which it does not respond to any packet.
 - If it receives a **DATA** packet, it broadcasts a **RERR** and resets the waiting timer (lifetime) to expire after the **current time PLUS delete_period**.
 - By the time the reboot node comes out of waiting and becomes active, **none of its neighbors will still be using it as an active next hop**.
 - Its own Seq# is updated once it receives a RREQ from any other node.